

R/openMP binding

F. Jamitzky

Leibniz Supercomputing Centre, Garching

jamitzky@lrz.de

ROMP



- **R openMP API**
- **R Syntax to Fortran Converter**
- **Accelerate R code by compilation**
- **Parallelize R code by vectorization**
- **Speedup by Compilation: ~100**
- **Speedup by Vectorization: ~100**
- **Total Speedup: ~10000**

Why R?

- **Very high abstraction level**
- **Lisp roots – “code that writes code”**
- **Interactivity – “Instant gratification”**
- **Fast prototyping language**
- **Huge Libraries – “Batteries included”**
- **Graphics and Plots – “nice and shiny”**

Why Fortran?



- **Well suited for numerical programming (very fast)**
- **Array arithmetics (syntax similar to R)**
- **Excelent R bindings (parts of R are written in Fortran)**

Why openMP?

- **Abstraction for vector processing**
- **Excelent Fortran bindings (Fortran and C are reference languages)**
- **Standard in high performance computing**
- **Excelent implementations, Fortran/openMP compiler from GNU, Intel, IBM, NAG, Microsoft, generating code for many CPUs and OSs.**

Philosophy



- **Use functional programming style**
- **Use closures**
- **R functions to Fortran functions in the “contains” part.**
- **Higher order functions: map/reduce**
- **Translate map/reduce to openMP for/reduce pragmas**

Abstractions



- R functions are translated to “pure” functions in Fortran
- R “**sum**” is replaced by “**sum.mp**”
- R “**apply**” is replaced by “**apply.mp**”
- Typing required, allowed types: int, double

Example

- **Compute distance of two time series:**

```
x = as.double(runif(100))
y = as.double(runif(100))
for(i in 1:100) res=res+(x[i]-y[i])**2
```

- **ROMP calls:**

```
sum.mp(dosum, (x[i]-y[i])**2, dbl(), i=1:100)
dosum.f = compile.mp(dosum(),
dbl(), x=dbl(100), y=dbl(100))
dosum.f(res=res, x=x, y=y)
```


Non-trivial Example: Pointwise Fractal Dimension



Compute pointwise dimension
of a cloud of points

Let N be the density of points at location \mathbf{x}

$$N(\mathbf{x}_i, r) = \sum_j \Theta(r - |\mathbf{x}_j - \mathbf{x}_i|).$$

where each point is smoothed with radius r

The fractal pointwise dimension is then defined as:

$$\alpha_i = (\log N(\mathbf{x}_i, r_2) - \log N(\mathbf{x}_i, r_1)) / (\log r_2 - \log r_1)$$

Ref: Local Scaling Properties for Diagnostic Purposes by W. Bunk, F. Jamitzky,
R. Pompl, C. Rath and G. Morfill, Springer 2002

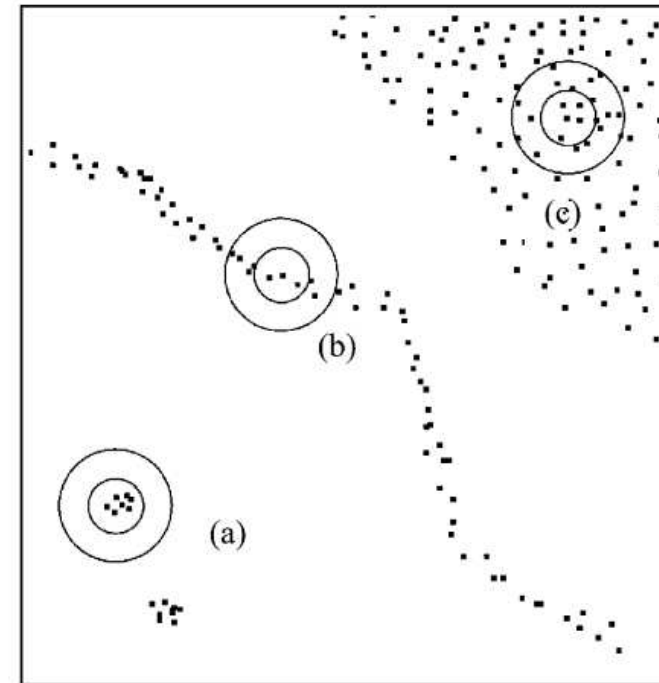


Fig. 1. Scheme illustrating the different dimensionality of point distributions. (a) A point-like structure. (b) A line-like structure. (c) An area-like structure.

Pure R style (verbose)



■ Compute local density of point set:

```
dist =  
  function(i,j,x,r)  
    ifelse(sum((x[i,1:ndim]-x[j,1:ndim])**2)>r**2,0,1)
```

```
dens_one =  
  function(j,x,r)  
    sum(sapply(1:np, function(i) dist(i,j,x,r)))
```

```
comp.dens =  
  function(x,r)  
    sapply(1:np, function(j) dens_one(j,x,r))
```

```
comp.dens(x, r=0.1)
```

“ROMP in style”



- **Compute local density of point set:**

```
sum.mp(dens_one,  
ifelse(sum((x[i,1:ndim]-x[j,1:ndim])**2)>r**2,0,1),  
int(), i=1:np, j=int())
```

```
apply.mp(dens, dens_one(j), int(np), j=1:np)
```

```
comp.dens = compile.mp( dens(),  
int(np),x=dbl(np,ndim),r=dbl(),ndim=int(),np=int())
```

```
comp.dens(x, r=0.1, ndim=3, np=100000)
```

Benchmarks




- **np=10000**
- **Pure R: time = 21800s = 6h!!**
- **ROMP: nproc=1 time = 3.2s**
- **ROMP: nproc=8 time = 0.6s**

Acceleration factor: >30000 !!

Benchmarks ROMP



- ROMP on HLRBII
- $np=10000$
- $nproc < 32$
- **Bad scaling**



Rmpi




- Rmpi: spawn R interpreter on each core
- applyLB MPI with load balancing

```
library(Rmpi)
mpi.bcast.Robj2slave(x)
mpi.applyLB(1:np,
  function(i)
    sum(sapply(1:np,
      function(j)
        ifelse(sum((x[i,1:ndim]-x[j,1:ndim])**2)>r**2,0,1)
      )))
```

Benchmark Rmpi



- Rmpi on HLRBII
- $np=10000$
- $nproc < 64$
- Strong scaling



01.08.08



Summary and Outlook

- **ROMP scales up to ~100 cores (SMP)**
- **Acceleration up to several 10000**
- **Pre Alpha Version**
- **Combination Rmpi+ROMP?**
- **Extending map/reduce: Use monads?**
- **Type inference aka automatic typing?**